



Introduction to virtualization and KVM

Contents

Origins

Objectives : pros and cons

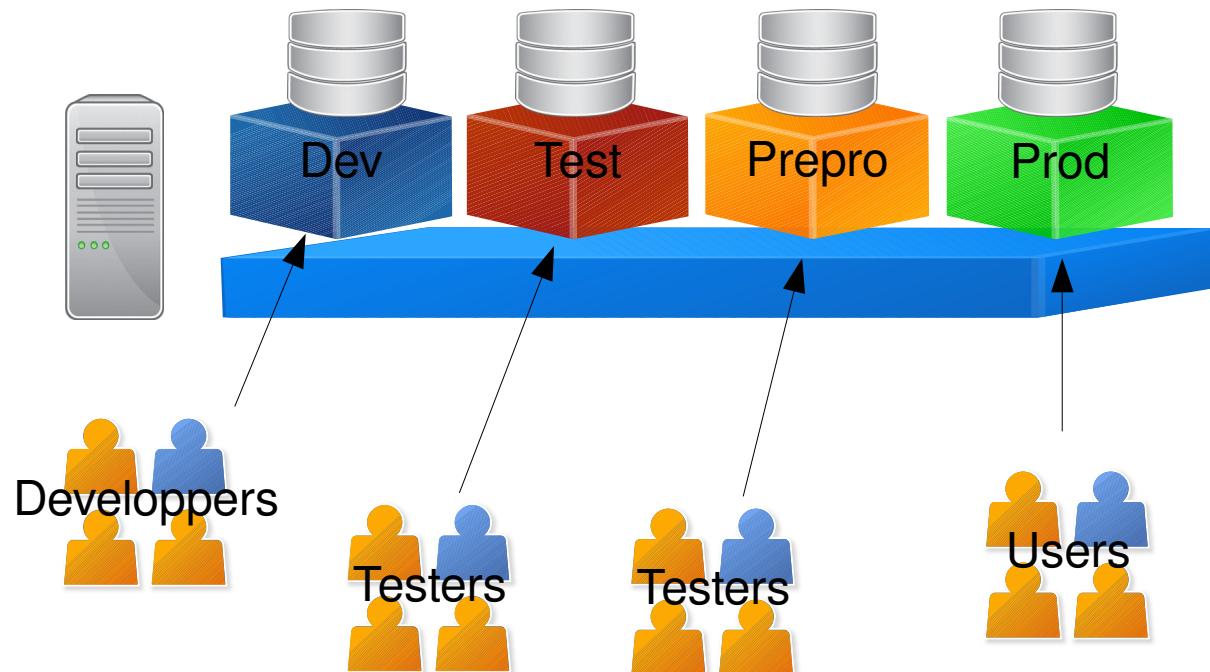
Technologies and examples

KVM

Origins

Mainframes

- In the 60's up to 80's
 - Mainframes
 - Only one (very expensive) computer, many environments
 - A solution :
 - environments are isolated
 - Their own adjustable resources



- In the 90's and 2000's
 - Virtualization on x86
 - Software emulation
 - Hardware virtualization (cpu instruction set)
 - Available on low cost hardware
 - A lot of free and non-free are available
 - Xen, KVM, Bochs, QEMU, Linux Vserver, VirtualBox
 - VirtualPC, VirtualServer, VMWare Server

Objectives

Objectives

- OS or application as a simple software
 - In a virtual environment (VE)
 - As many computers and many OS on a single computer
- Tests
 - Isolated VE
 - Migrate VE to more adapted hardware
- Principles
 - A host OS
 - A virtualisation software : hypervisor
 - Other OS : guests

Cost reduction

- energy consumption
 - servers
 - air conditioning
 - Green factor
- hardware maintenance delays
 - Reduce maintenance procedures
 - ... the administration time consumption
- Reduce hardware compatibility management
 - few components expenditures
 - few suppliers to manage
 - Many versions of software can run longer
- Hardware used efficiently
 - high cost components (CPU, RAM,...) are used more or less full time
 - Consolidation
 - Less hardware
 - Less room space

- Reduce maintenance procedures
 - Install from predefined models
 - With automated IT infrastructure integration (network, backup agent, supervision agent, ...)
- Central monitoring
 - A central supervision point
 - To decide migration action
 - To react
- Start/stop VM without restart physical server
 - Snapshots : quick restore

Disavantages

Limitations

- **Compatibility**

- Software may not be compatible with virtualization
- Then exceptions to manage

- **Security**

- Hardware = single point of failure
- Data protection with delegated data center
- Server sharing

- **VM out of control**

- Too much VM where fewer are needed
- Orphan VM

Technologies

- Host OS isolates user spaces
 - One user space = one execution environment
 - No guest OS
 - Low overhead rate = high performances
 - Host OS dependant
- Examples
 - Chroot
 - Docker
 - Linux Vserver
 - LXC

Full virtualization

- Hypervisor first generation
 - Host OS with light kernel
 - Optimize transfers between guests OS and hardware
- Guest OS adaptation ?
 - No
 - If yes, paravirtualization (guests and host OS's cooperation)
- High performances because thin layer
- Examples
 - Xen
 - Vsphere
 - Hyper-V server

- Hypervisor second generation
 - Close to emulation hardware components
 - Good isolation
 - Thicker layer between hardware and guests OS
 - Performances reduced
- Examples
 - VirtualBox
 - KVM
 - QEMU
 - Bochs

Kernel-based Virtual Machine (KVM)

Introduction

Preparation

Guest creation

VMs management

Introduction

What is KVM ?

- Hypervisor in a kernel module

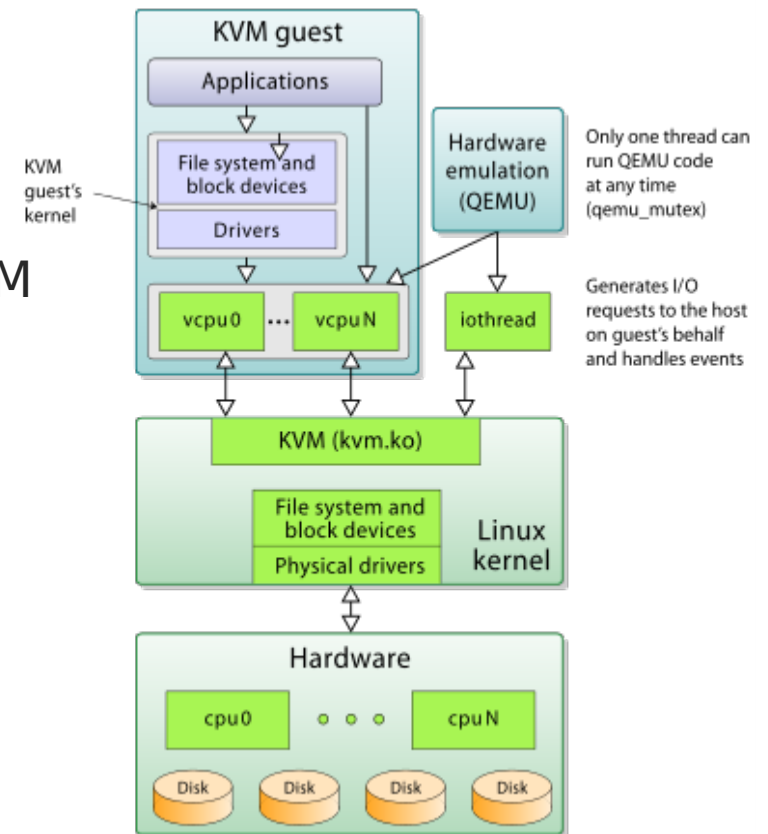
- In Linux mainline since 2007
- Also on FreeBSD
- Architectures : x86, S/390, PowerPC, ARM

- Guests

- Linux, BSD, Solaris
- Windows
- OS X


- I/O hardware emulation

- Based on QEMU
- In user- space



```
$ uname -r
3.16.0-4-amd64
$ modinfo kvm
filename:      /lib/modules/3.16.0-4-amd64/kernel/arch/x86/kvm/kvm.ko
license:      GPL
author:       Qumranet
depends:
intree:       Y
vermagic:     3.16.0-4-amd64 SMP mod_unload modversions
parm:         ignore_msrs:bool
parm:         min_timer_period_us:uint
parm:         tsc_tolerance_ppm:uint
```

- Open source
 - Since 2005 (Qumranet company)
 - Since 2008, RedHat
 - 2012 : v 1.2
- Processors with Virtualization Technology
 - Check BIOS
 - Check CPU (/proc/cpuinfo)
- Libvirt (*virsh*, *virt-manager*)
 - Independent and standard VM manager on a host
 - API to provision, create, modify, monitor, start, stop, migrate VM



- Virtualized and emulated

- KVM emulates « full software » devices
- Virtual CPUs
- PCI bridges, USB controllers, IDE, AHCI controller, etc.
- Graphics, sound, network (**e1000**, **rtl8139**)

- Para-virtualized devices

- Faster I/O transfers, higher I/O throughput
- Based on kernel modules : **virtio-***
- Network, disks, SCSI, etc

- Physical devices

- Direct access to certain hardware : passthrough
- Depends on platforms

- **Storage pools**
 - Local SP : attached to the host server (local devices)
 - Network shared SP
- **Emulated storage devices**
 - Virtio-scsi : recommended
 - IDE, CDRROM, USB mass storage, AHCI for SATA
- **Host storage**
 - Raw for I/O performances
 - QCOW2 for advanced features : snapshots, encryption, compression

Preparation

For about 10 VM :

- CPU

- At least 4 cores
- Virtualization technology

- Memory

- At least 8 GB

- Disk drive

- At least 100 GB

- *Overcommitting technology*

- it is possible to allocate more virtualized CPU/RAM than physical
- KVM distributes CPU/RAM to VM's

- Qemu-kvm
 - Qemu binaries
 - User-level emulator
- Libvirt
 - GUI tool
 - Virtualization management tools
 - Start
- Guestfs
 - Tool for offline images management

- According to VM guest goals, choose :
 - number of CPU
 - RAM amount
 - Storage volume/technology and select host disk
 - Networking
 - Bridge (check bridge is set on host)
 - NAT
 - fixed or DHCP
 - OS installation method : medium, NFS, HTTP, copy from pattern, unattended

Guest creation

- Online command **virt-install**
 - **virt-install **
**--network bridge:br0 **
**--name rh7dev **
**--ram=1024 **
**--vcpus=1 **
**--disk path=/opt/vm-images/rh7dev.img,size=10 **
--cdrom /opt/iso/rhel7.3.iso
 - Other interesting options :
 - location** : URL of the distribution tree installation
 - os-type** : guest operating system
 - arch** : if non-native architecture
- The disk is named **vda**.

- Cloning is another way to create a guest

- Stop all I/O on the VM to be cloned :

```
virsh suspend rhel7dev
```

- Run the cloning command :

```
virt-clone \  
    --connect qemu:///system \  
    --original rhel7dev \  
    --name rhel7dev-clone \  
    --file /opt/vm-images/rhel7dev-clone.img
```

- Resume original VM :

```
virtsh resume rhel7dev
```

- Start the cloned VM

```
virsh start rhel7dev
```

VM management

Current tasks

- List VMs
 - `virtsh list --all`
- Show VM information
 - `virtsh dominfo rhel7dev`
- Show VCPU/Memory usage
 - `virt-top`
- Show VM disk partitions
 - `virt-df rhel7dev`
- Start VM
 - `virsh start rhel7dev`
- Stop VM (shutdown the OS)
 - `virsh shutdown rhel7dev`

- Access to VM console

- **virsh shutdown rhel7dev**

first add « **console=tty0 console=ttyS0,115200** » to kernel boot line in grub configuration of VM

- Attach/detach storage device

- **virsh attach-disk rhel7dev /dev/sdb vdb **
--driver qemu --mode shareable

vdb : device name mapped in th VM

- **virsh detach-disk rhel7dev vdb**

Changing VM memory dynamically

- **Change dynamically memory**

- View current's memory settings

```
virsh dominfo rhel7dev | grep memory
```

```
Max memory : 1048576 kB
```

```
Used memory : 1048576 kB
```

- Set to 512 MB (value in kB)

```
virsh setmem rhel7dev 524288
```

```
virsh dominfo rhel7dev | grep memory
```

```
Max memory : 1048576 kB
```

```
Used memory : 524288 kB
```


- Increase memory

- Stop VM

```
virsh shutdown rhel7dev
```

- Edit VM's configuration file

```
virsh edit rhel7dev
```

```
<memory unit='KiB'>2097152</memory>
```

- Restart the VM with the new parameter

```
virsh create /etc/libvirt/qemu/rhel7dev.xml
```

- Check memory settings

```
virsh dominfo vm1 | grep memory
```

```
Max memory:      2097152 kB
```

```
Used memory:     524288 kB
```

Change VPCUs

- Stop the VM

```
virsh shutdown rhel7dev
```

- Modify the CPU in the configuration file

```
virsh edit rhel7dev
```

```
<vcpu>1</vcpu>
```

- Restart VM

```
virsh create /etc/libvirt/qemu/rhel7dev.xml
```

Add disk

- Create a 10 GB file

```
dd if=/dev/zero of=/opt/vm-images/rhel7dev-hd2.img  
bs=1M count=10240
```

- Shutdown the VM

```
virsh shutdown rhel7dev
```

- Add a disk stanza in the VM configuration file

```
<disk type='file' device='disk'>  
  <driver name='qemu' type='raw' cache='none' io='threads' />  
  <source file='/opt/vm-images/rhel7dev-hd2.img' />  
  <target dev='vdb' bus='virtio' />  
  <address type='pci' domain='0x0000' bus='0x00' slot='0x06'  
    function='0x0' />  
</disk>
```

- Restart

```
virsh create /etc/libvirt/qemu/rhel7dev.xml
```

Deleting VMs

- Shutdown the VM

```
virsh shutdown rhel7dev
```

- If the VM is not responding or fails to shutdown, shut it down forcefully

```
virsh destroy rhel7dev
```

- Undefine the VMs configuration

```
virsh undefine rhel7dev
```

- Remove the VM's image files

```
rm /opt/vm-images/rhel7dev*.img
```

Improve I/O performances

- Use VirtIO drivers (disk, network)

```
<disk type='file' device='disk'>  
  <source file='/opt/vm-images/disk1.img' />  
  <target dev='vda' bus='virtio' />  
</disk>
```

- Assign device disk instead of « file » disk

```
<source file='/dev/vgvm/lvdd1' />
```

Thanks for your attention

LINAGORA - headquarters

80, rue Roque de Fillol
92800 PUTEAUX
FRANCE

Phone : +33 (0)1 46 96 63 63

Info : info@linagora.com

Web : www.linagora.com

 **@linagora**

 **facebook.com/Linagora/**

